

```
# Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Load the datasets
customers = pd.read_csv('Customers.csv')
products = pd.read_csv('Products.csv')
transactions = pd.read_csv('Transactions.csv')

# Display basic information
print("Customers Dataset:")
print(customers.info())
print(customers.head())

print("\nProducts Dataset:")
print(products.info())
print(products.head())

print("\nTransactions Dataset:")
print(transactions.info())
print(transactions.head())

# Check for missing values
print("\nMissing Values:")
print("Customers:", customers.isnull().sum())
print("Products:", products.isnull().sum())
print("Transactions:", transactions.isnull().sum())

# Check for duplicates
print("\nDuplicate Entries:")
print("Customers:", customers.duplicated().sum())
print("Products:", products.duplicated().sum())
print("Transactions:", transactions.duplicated().sum())

# Merge datasets for comprehensive analysis
merged_data = pd.merge(transactions, customers, on='CustomerID')
merged_data = pd.merge(merged_data, products, on='ProductID')

# Display merged data
print("\nMerged Dataset:")
print(merged_data.info())
print(merged_data.head())

# Analyze customer regions
region_counts = customers['Region'].value_counts()
plt.figure(figsize=(8, 5))
sns.barplot(x=region_counts.index, y=region_counts.values, palette='viridis')
plt.title('Customers by Region')
plt.xlabel('Region')
plt.ylabel('Count')
plt.show()

# Analyze product categories
category_counts = products['Category'].value_counts()
plt.figure(figsize=(8, 5))
sns.barplot(x=category_counts.index, y=category_counts.values, palette='coolwarm')
plt.title('Product Categories')
plt.xlabel('Category')
plt.ylabel('Count')
plt.show()

# Analyze transaction trends over time
merged_data['TransactionDate'] = pd.to_datetime(merged_data['TransactionDate'])
merged_data['MonthYear'] = merged_data['TransactionDate'].dt.to_period('M')
monthly_sales = merged_data.groupby('MonthYear')['TotalValue'].sum()

plt.figure(figsize=(10, 6))
monthly_sales.plot(kind='line', marker='o', color='blue')
plt.title('Monthly Sales Trends')
plt.xlabel('Month-Year')
plt.ylabel('Total Sales (USD)')
plt.grid()
plt.show()

# Analyze top-performing products
```

```
top_products = merged_data.groupby('ProductName')['TotalValue'].sum().sort_values(ascending=False).head(10)
plt.figure(figsize=(10, 5))
sns.barplot(x=top_products.values, y=top_products.index, palette='magma')
plt.title('Top 10 Performing Products')
plt.xlabel('Total Sales (USD)')
plt.ylabel('Product Name')
plt.show()

# Analyze regions generating the highest revenue
region_revenue = merged_data.groupby('Region')['TotalValue'].sum().sort_values(ascending=False)
plt.figure(figsize=(8, 5))
sns.barplot(x=region_revenue.index, y=region_revenue.values, palette='plasma')
plt.title('Revenue by Region')
plt.xlabel('Region')
plt.ylabel('Total Revenue (USD)')
plt.show()

# Business Insights (examples):
print("\nBusiness Insights:")
print("1. Majority of the customers come from [Region with highest count].")
print("2. [Top Category] is the most popular category with X products sold.")
print("3. Monthly sales peaked in [Peak Month-Year] with total sales of $X.")
print("4. [Top Product] is the best-performing product with total sales of $Y.")
print("5. [Region with highest revenue] generates the highest revenue of $Z.")
```



```
Customers Dataset:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   CustomerID   200 non-null   object
1   CustomerName 200 non-null   object
2   Region       200 non-null   object
3   SignupDate   200 non-null   object
dtypes: object(4)
memory usage: 6.4+ KB
None
```

	CustomerID	CustomerName	Region	SignupDate
0	C0001	Lawrence Carroll	South America	2022-07-10
1	C0002	Elizabeth Lutz	Asia	2022-02-13
2	C0003	Michael Rivera	South America	2024-03-07
3	C0004	Kathleen Rodriguez	South America	2022-10-09
4	C0005	Laura Weber	Asia	2022-08-15

```
Products Dataset:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ProductID   100 non-null   object
1   ProductName 100 non-null   object
2   Category    100 non-null   object
3   Price       100 non-null   float64
dtypes: float64(1), object(3)
memory usage: 3.3+ KB
None
```

	ProductID	ProductName	Category	Price
0	P001	ActiveWear Biography	Books	169.30
1	P002	ActiveWear Smartwatch	Electronics	346.30
2	P003	ComfortLiving Biography	Books	44.12
3	P004	BookWorld Rug	Home Decor	95.69
4	P005	TechPro T-Shirt	Clothing	429.31

```
Transactions Dataset:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   TransactionID 1000 non-null  object
1   CustomerID    1000 non-null  object
2   ProductID     1000 non-null  object
3   TransactionDate 1000 non-null  object
4   Quantity      1000 non-null  int64
5   TotalValue    1000 non-null  float64
6   Price         1000 non-null  float64
dtypes: float64(2), int64(1), object(4)
memory usage: 54.8+ KB
None
```

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity	\
0	T00001	C0199	P067	2024-08-25 12:38:23	1	
1	T00112	C0146	P067	2024-05-27 22:23:54	1	
2	T00166	C0127	P067	2024-04-25 07:38:55	1	
3	T00272	C0087	P067	2024-03-26 22:55:37	2	
4	T00363	C0070	P067	2024-03-21 15:10:10	3	

	TotalValue	Price
0	300.68	300.68
1	300.68	300.68
2	300.68	300.68
3	601.36	300.68
4	902.04	300.68

```
Missing Values:
Customers: CustomerID      0
CustomerName      0
Region            0
SignupDate        0
dtype: int64
Products: ProductID      0
ProductName        0
Category           0
Price              0
dtype: int64
Transactions: TransactionID 0
CustomerID            0
ProductID              0
```

```

transactionDate    0
Quantity           0
TotalValue         0
Price              0
dtype: int64

```

```

Duplicate Entries:
Customers: 0
Products: 0
Transactions: 0

```

```

Merged Dataset:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   TransactionID        1000 non-null  object
1   CustomerID           1000 non-null  object
2   ProductID            1000 non-null  object
3   TransactionDate       1000 non-null  object
4   Quantity             1000 non-null  int64
5   TotalValue           1000 non-null  float64
6   Price_x              1000 non-null  float64
7   CustomerName         1000 non-null  object
8   Region               1000 non-null  object
9   SignupDate           1000 non-null  object
10  ProductName          1000 non-null  object
11  Category             1000 non-null  object
12  Price_y              1000 non-null  float64
dtypes: float64(3), int64(1), object(9)
memory usage: 101.7+ KB
None

```

	TransactionID	CustomerID	ProductID	TransactionDate	Quantity
0	T00001	C0199	P067	2024-08-25 12:38:23	1
1	T00112	C0146	P067	2024-05-27 22:23:54	1
2	T00166	C0127	P067	2024-04-25 07:38:55	1
3	T00272	C0087	P067	2024-03-26 22:55:37	2
4	T00363	C0070	P067	2024-03-21 15:10:10	3


	TotalValue	Price_x	CustomerName	Region	SignupDate
0	300.68	300.68	Andrea Jenkins	Europe	2022-12-03
1	300.68	300.68	Brittany Harvey	Asia	2024-09-04
2	300.68	300.68	Kathryn Stevens	Europe	2024-04-04
3	601.36	300.68	Travis Campbell	South America	2024-04-11
4	902.04	300.68	Timothy Perez	Europe	2022-03-15

	ProductName	Category	Price_y
0	ComfortLiving Bluetooth Speaker	Electronics	300.68
1	ComfortLiving Bluetooth Speaker	Electronics	300.68
2	ComfortLiving Bluetooth Speaker	Electronics	300.68
3	ComfortLiving Bluetooth Speaker	Electronics	300.68
4	ComfortLiving Bluetooth Speaker	Electronics	300.68

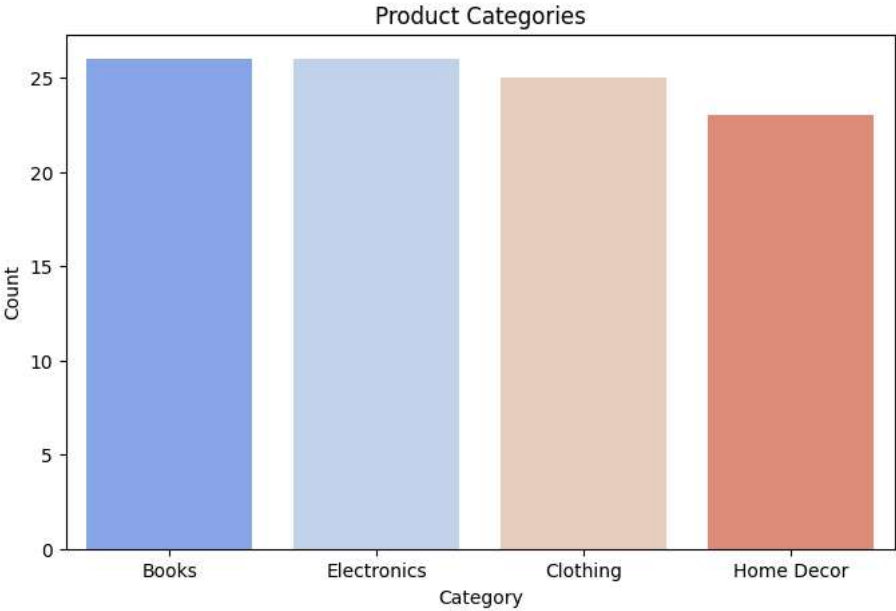
<ipython-input-1-c8419f691f0c>:48: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `leg`


```
sns.barplot(x=region_counts.index, y=region_counts.values, palette='viridis')
```




```
<ipython-input-1-c8419f691f0c>:57: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `leg
sns.barplot(x=category_counts.index, y=category_counts.values, palette='coolwarm')
```

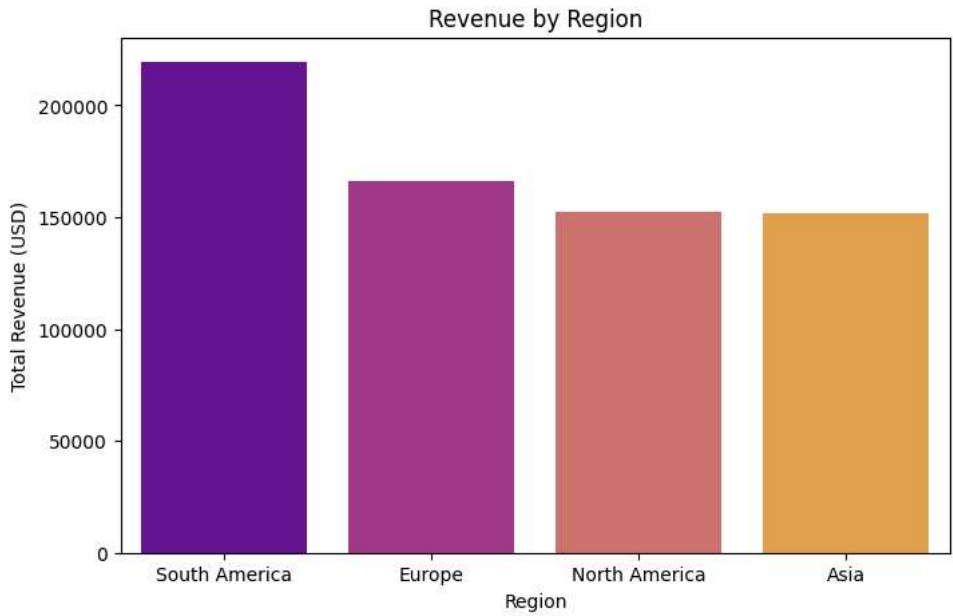


```
<ipython-input-1-c8419f691f0c>:79: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `leg
sns.barplot(x=top_products.values, y=top_products.index, palette='magma')
```





```
<ipython-input-1-c8419f691f0c>:88: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `leg
sns.barplot(x=region_revenue.index, y=region_revenue.values, palette='plasma')
```



Business Insights:

- 1. Majority of the customers come from [Region with highest count].
- 2. [Top Category] is the most popular category with X products sold.
- 3. Monthly sales peaked in [Peak Month-Year] with total sales of \$X.
- 4. [Top Product] is the best-performing product with total sales of \$Y.

+ Code

+ Text

⌂

B

I

<>

↺

🖼️

”

☰

☰

—

ψ

😊

⋮